



# FRUIT CLASSIFICATION MODEL

## ASSIGNMENT TWO

### GROUP MEMBERS

#### **STEP ONE (IMPORTING LIBRARIES, DATASET & INSPECTING THE DATASET**

- ARINJUNA SARAH 2020/BCS/093/PS

#### **STEP TWO (PREPROCESSING DATA & SPLITTING THE DATA**

- OKWAKUNDA GLORIA 2020/BCS/091/PS
- KAYINZA MARIAM 2020/BCS/036/PS

#### **STEP THREE ( DATA STANDARDIZATION, ARGUMENTATION, CACHING & SHUFFLING)**

- AINAMANI CHRISTIAN 2020/BCS/001
- MAYANJA ROBERT 2020/BCS/043/PS

#### **STEP FOUR (Building the Model, ARCHITECTURE & TRAINING)**

- WESONGA BOB 2020/BCS/092/PS
- SIEMBA ERNEST OOKO 2020/BCS/005

#### **STEP FIVE (EVALUATION & EXPORTING THE MODEL**

- ATWANZIRE TIMOTHY IAN 2020/BCS/026/PS

- KATUSHABE MOREEN 2020/BCS/034/PS

<https://www.kaggle.com/datasets/arnavmehta710a/fids30> (<https://www.kaggle.com/datasets/arnavmehta710a/fids30>) OR <https://www.vicos.si/resources/fids30> (<https://www.vicos.si/resources/fids30>)

## Importing Libraries

**ARINJUNA SARAH**

```
In [1]: import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt
from tensorflow import keras
import matplotlib.pyplot as plt
import numpy as np
import os
```

2023-04-03 10:53:31.945909: I tensorflow/core/platform/cpu\_feature\_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2  
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

## Set of constants

- Batch Size - A number of samples processed before the model is updated
- Image Size - The Resolution
- Channels - Number of steps of the shifting of a convolutional filter over an input image
- Epochs - Total number of iterations of all the training data in one cycle for training the machine learning model
- Data\_dir - Path of the dataset

```
In [2]: BATCH_SIZE = 32
IMAGE_SIZE = 256
CHANNELS=30
EPOCHS=30
data_dir = "FIDS30"
```

## Importing the Dataset

- We will <https://www.vicos.si/resources/fids30> (<https://www.vicos.si/resources/fids30>) data which can also be got from <https://www.kaggle.com/datasets/arnavmehta710a/fids30> (<https://www.kaggle.com/datasets/arnavmehta710a/fids30>)
- We will use the "tf.keras.preprocessing.image\_dataset\_from\_directory" inbuild func to load

the dataset

- It takes in
  - "seeds" @para for defining a starting point for the sequence
  - "shuffle" @para to aim to mix up data and can optionally retain logical relationships between columns
  - image\_size && batch\_size

```
In [3]: dataset = tf.keras.preprocessing.image_dataset_from_directory(
        data_dir,
        seed=123,
        shuffle=True,
        image_size=(IMAGE_SIZE, IMAGE_SIZE),
        batch_size=BATCH_SIZE
    )
```

Found 971 files belonging to 30 classes.

2023-04-03 10:53:36.872155: I tensorflow/core/platform/cpu\_feature\_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: SSE4.1 SSE4.2

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

2023-04-03 10:53:36.874447: I tensorflow/core/common\_runtime/process\_util.cc:146] Creating new thread pool with default inter op setting: 2. Tune using inter\_op\_parallelism\_threads for best performance.

**ClassNames - These are classes of our dataset**

```
In [4]: class_names = dataset.class_names
print(len(class_names))
class_names
```

30

```
Out[4]: ['acerolas',
         'apples',
         'apricots',
         'avocados',
         'bananas',
         'blackberries',
         'blueberries',
         'cantaloupes',
         'cherries',
         'coconuts',
         'figs',
         'grapefruits',
         'grapes',
         'guava',
         'kiwifruit',
         'lemons',
         'limes',
         'mangos',
         'olives',
         'oranges',
         'passionfruit',
         'peaches',
         'pears',
         'pineapples',
         'plums',
         'pomegranates',
         'raspberries',
         'strawberries',
         'tomatoes',
         'watermelons']
```

## Preprocess the Dataset

*Kayinza Mariam, Okwakunda Gloria & Katushabe Moreen*

### The shape of the first batch of images in a dataset

- As you can see above, each element in the dataset is a tuple. First element is a batch of 32 elements of images. Second element is a batch of 32 elements of class labels

```
In [5]: for image_batch, labels_batch in dataset.take(1):
        print(image_batch.shape)
        print(labels_batch.numpy())
```

```
(32, 256, 256, 3)
[ 5 11 27 21 24 19  3 22  6 14 11 19 17 24 11 21 22 12 21 17  4 22  5
25
 1  8 25  6  4 21 23 18]
```

## Visualizing the Images

```
In [6]: plt.figure(figsize=(10, 10))
        for image_batch, labels_batch in dataset.take(1):
            for i in range(12):
                ax = plt.subplot(3, 4, i + 1)
                plt.imshow(image_batch[i].numpy().astype("uint8"))
                plt.title(class_names[labels_batch[i]])
                plt.axis("off")
```

pomegranates



blackberries



apricots



pomegranates




avocados



oranges



oranges



peaches



cherries



pomegranates



acerolas



tomatoes



## Splitting the dataset

- It's good practice to use a validation split when developing your model.
- We will use 80% of the images for training and 20% for validation.

```
In [7]: len(dataset)
```

```
Out[7]: 31
```

```
In [8]: train_size = 0.8  
len(dataset)*train_size
```

```
Out[8]: 24.8
```

```
In [9]: train_ds = dataset.take(24)  
len(train_ds)
```

```
Out[9]: 24
```

```
In [10]: test_ds = dataset.skip(24)  
len(test_ds)
```

```
Out[10]: 7
```

```
In [11]: val_size=0.1  
len(dataset)*val_size
```

```
Out[11]: 3.1
```

```
In [12]: val_ds = test_ds.take(3)  
len(val_ds)
```

```
Out[12]: 3
```

```
In [13]: test_ds = test_ds.skip(3)  
len(test_ds)
```

```
Out[13]: 4
```

```
In [14]: def get_dataset_partitions_tf(ds, train_split=0.8, val_split=0.1, test_
         assert (train_split + test_split + val_split) == 1

         ds_size = len(ds)

         if shuffle:
             ds = ds.shuffle(shuffle_size, seed=12)

         train_size = int(train_split * ds_size)
         val_size = int(val_split * ds_size)

         train_ds = ds.take(train_size)
         val_ds = ds.skip(train_size).take(val_size)
         test_ds = ds.skip(train_size).skip(val_size)

         return train_ds, val_ds, test_ds
```

```
In [15]: train_ds, val_ds, test_ds = get_dataset_partitions_tf(dataset)
```

```
In [16]: len(train_ds)
```

```
Out[16]: 24
```

```
In [17]: len(val_ds)
```

```
Out[17]: 3
```

```
In [18]: len(test_ds)
```

```
Out[18]: 4
```

## DATA STANDARDIZATION, ARGUMENTATION, CACHING & SHUFFLING

*Ainamani Christian & MAYANJA ROBERT*

### Cache, Shuffle, and Prefetch the Dataset

- "Cache" a module for language modelling which stores previous hidden states in memory cells
- "Shuffle" a module to mixing up dataset
- "Prefetch" overlaps the preprocessing and model execution of a training step.

```
In [19]: train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.
         val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTO
         test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AU
```

## Creating a Layer for Resizing and Normalization

```
In [20]: resize_and_rescale = tf.keras.Sequential([
          layers.experimental.preprocessing.Resizing(IMAGE_SIZE, IMAGE_SIZE),
          layers.experimental.preprocessing.Rescaling(1./255),
          ])
```

## Data Augmentation

```
In [21]: data_augmentation = tf.keras.Sequential([
          layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),
          layers.experimental.preprocessing.RandomRotation(0.2),
          ])
```

## Applying Data Augmentation to Train Dataset

```
In [22]: train_ds = train_ds.map(
          lambda x, y: (data_augmentation(x, training=True), y)
          ).prefetch(buffer_size=tf.data.AUTOTUNE)
```

WARNING:tensorflow:From /home/fedora/anaconda3/lib/python3.10/site-packages/tensorflow/python/autograph/pyct/static\_analysis/liveness.py:83: Analyzer.lamba\_check (from tensorflow.python.autograph.pyct.static\_analysis.liveness) is deprecated and will be removed after 2023-09-23.

Instructions for updating:

Lambda fuctions will be no more assumed to be used in the statement where they are used, or at least in the same block. <https://github.com/tensorflow/tensorflow/issues/56089> (<https://github.com/tensorflow/tensorflow/issues/56089>)

WARNING:tensorflow:Using a while\_loop for converting RngReadAndSkip cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting ImageProjectiveTransformV3 cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting RngReadAndSkip cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting Bitcast cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting StatelessRandomUniformV2 cause there is no registered converter for this op.

WARNING:tensorflow:Using a while\_loop for converting ImageProjectiveTransformV3 cause there is no registered converter for this op.



# Building the Model

*Wesonga Bob & Siemba Ernest Ooko*

## Model Architecture

```
In [23]: num_classes = 30

model = tf.keras.Sequential([
    tf.keras.layers.Rescaling(1./255),
    tf.keras.layers.Conv2D(32, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(32, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(32, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(num_classes)
])
```

```
In [24]: model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])
```

```
In [25]: history = model.fit(
          train_ds,
          validation_data=val_ds,
          epochs=EPOCHS
        )
```

Epoch 1/30

```
2023-04-03 10:53:54.937934: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 10:54:01.410207: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:392] Filling up shuffle buffer (this may take a while): 17
of 10000
2023-04-03 10:54:01.412893: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 10:54:02.864622: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 10:54:03.299995: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:417] Shuffle buffer filled.
2023-04-03 10:54:03.300049: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:392] Filling up shuffle buffer (this may take a while): 1 o
f 1000
2023-04-03 10:54:03.300140: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:417] Shuffle buffer filled.
```

24/24 [=====] - ETA: 0s - loss: 3.3086 - accuracy: 0.0768

```
2023-04-03 10:55:37.312737: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 10:55:42.202548: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 10:55:43.407733: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 10:55:43.667746: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:392] Filling up shuffle buffer (this may take a while): 23
of 10000
2023-04-03 10:55:43.744196: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:417] Shuffle buffer filled.
2023-04-03 10:55:43.744436: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:392] Filling up shuffle buffer (this may take a while): 1 o
f 1000
2023-04-03 10:55:43.744552: I tensorflow/core/kernels/data/shuffle_data
set_op.cc:417] Shuffle buffer filled.
```

24/24 [=====] - 117s 4s/step - loss: 3.3086 - accuracy: 0.0768 - val\_loss: 3.1436 - val\_accuracy: 0.1042

Epoch 2/30

24/24 [=====] - 91s 4s/step - loss: 2.7993 - accuracy: 0.1706 - val\_loss: 2.5101 - val\_accuracy: 0.2083

Epoch 3/30

24/24 [=====] - 90s 4s/step - loss: 2.3504 - accuracy: 0.2461 - val\_loss: 2.6380 - val\_accuracy: 0.1771

Epoch 4/30

24/24 [=====] - 92s 4s/step - loss: 2.1902 - accuracy: 0.2878 - val\_loss: 2.3530 - val\_accuracy: 0.2188

```
Epoch 5/30
24/24 [=====] - 90s 4s/step - loss: 1.9695 -
accuracy: 0.3529 - val_loss: 2.3430 - val_accuracy: 0.2292
Epoch 6/30
24/24 [=====] - 91s 4s/step - loss: 1.8728 -
accuracy: 0.4102 - val_loss: 2.1488 - val_accuracy: 0.3438
Epoch 7/30
24/24 [=====] - 88s 4s/step - loss: 1.7951 -
accuracy: 0.4089 - val_loss: 2.2537 - val_accuracy: 0.3125
Epoch 8/30
24/24 [=====] - 89s 4s/step - loss: 1.5407 -
accuracy: 0.4935 - val_loss: 2.1685 - val_accuracy: 0.3021
Epoch 9/30
24/24 [=====] - 88s 4s/step - loss: 1.5102 -
accuracy: 0.4857 - val_loss: 2.0344 - val_accuracy: 0.4167
Epoch 10/30
24/24 [=====] - 90s 4s/step - loss: 1.3796 -
accuracy: 0.5495 - val_loss: 1.8846 - val_accuracy: 0.4375
Epoch 11/30
24/24 [=====] - 90s 4s/step - loss: 1.3137 -
accuracy: 0.5664 - val_loss: 1.7198 - val_accuracy: 0.5104
Epoch 12/30
24/24 [=====] - 89s 4s/step - loss: 1.1813 -
accuracy: 0.6263 - val_loss: 2.1887 - val_accuracy: 0.3854
Epoch 13/30
24/24 [=====] - 90s 4s/step - loss: 1.0463 -
accuracy: 0.6536 - val_loss: 2.2246 - val_accuracy: 0.3958
Epoch 14/30
24/24 [=====] - 89s 4s/step - loss: 1.0228 -
accuracy: 0.6693 - val_loss: 2.0088 - val_accuracy: 0.3646
Epoch 15/30
24/24 [=====] - 88s 4s/step - loss: 0.9698 -
accuracy: 0.6810 - val_loss: 1.8004 - val_accuracy: 0.4479
Epoch 16/30
24/24 [=====] - 89s 4s/step - loss: 0.9047 -
accuracy: 0.7148 - val_loss: 1.6077 - val_accuracy: 0.5208
Epoch 17/30
24/24 [=====] - 88s 4s/step - loss: 0.8362 -
accuracy: 0.7383 - val_loss: 1.6431 - val_accuracy: 0.5208
Epoch 18/30
24/24 [=====] - 90s 4s/step - loss: 0.7006 -
accuracy: 0.7656 - val_loss: 1.7465 - val_accuracy: 0.4896
Epoch 19/30
24/24 [=====] - 90s 4s/step - loss: 0.6939 -
accuracy: 0.7786 - val_loss: 1.7895 - val_accuracy: 0.4792
Epoch 20/30
24/24 [=====] - 88s 4s/step - loss: 0.6325 -
accuracy: 0.7747 - val_loss: 1.8770 - val_accuracy: 0.5312
Epoch 21/30
24/24 [=====] - 90s 4s/step - loss: 0.6489 -
accuracy: 0.7839 - val_loss: 1.9455 - val_accuracy: 0.5104
Epoch 22/30
24/24 [=====] - 89s 4s/step - loss: 0.5746 -
accuracy: 0.8164 - val_loss: 1.9166 - val_accuracy: 0.5312
Epoch 23/30
24/24 [=====] - 89s 4s/step - loss: 0.5663 -
```

```

accuracy: 0.8177 - val_loss: 1.7853 - val_accuracy: 0.5312
Epoch 24/30
24/24 [=====] - 90s 4s/step - loss: 0.5305 -
accuracy: 0.8125 - val_loss: 1.7998 - val_accuracy: 0.5000
Epoch 25/30
24/24 [=====] - 95s 4s/step - loss: 0.4274 -
accuracy: 0.8620 - val_loss: 1.9640 - val_accuracy: 0.5104
Epoch 26/30
24/24 [=====] - 112s 5s/step - loss: 0.4280 -
accuracy: 0.8620 - val_loss: 2.0358 - val_accuracy: 0.5208
Epoch 27/30
24/24 [=====] - 95s 4s/step - loss: 0.4456 -
accuracy: 0.8594 - val_loss: 1.7088 - val_accuracy: 0.5938
Epoch 28/30
24/24 [=====] - 90s 4s/step - loss: 0.3672 -
accuracy: 0.8841 - val_loss: 1.8127 - val_accuracy: 0.5208
Epoch 29/30
24/24 [=====] - 105s 4s/step - loss: 0.3988 -
accuracy: 0.8633 - val_loss: 2.0113 - val_accuracy: 0.5417
Epoch 30/30
24/24 [=====] - 91s 4s/step - loss: 0.3796 -

```

## Evaluation

In [26]: `scores = model.evaluate(test_ds)`

```

2023-04-03 11:41:02.114465: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 11:41:07.002962: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 11:41:08.190928: I tensorflow/core/kernels/data/shuffle_dat
aset_op.cc:392] Filling up shuffle buffer (this may take a while): 21
of 10000
2023-04-03 11:41:08.220648: W tensorflow/core/lib/png/png_io.cc:88] PNG
warning: iCCP: known incorrect sRGB profile
2023-04-03 11:41:08.521201: I tensorflow/core/kernels/data/shuffle_dat
aset_op.cc:417] Shuffle buffer filled.
2023-04-03 11:41:08.521285: I tensorflow/core/kernels/data/shuffle_dat
aset_op.cc:392] Filling up shuffle buffer (this may take a while): 1 o
f 1000
2023-04-03 11:41:08.521357: I tensorflow/core/kernels/data/shuffle_dat
aset_op.cc:417] Shuffle buffer filled.

4/4 [=====] - 15s 1s/step - loss: 2.1635 - ac
curacy: 0.5469

```

In [27]: `# Scores is just a list containing loss and accuracy value`  
`scores`

Out[27]: `[2.1634774208068848, 0.546875]`

## Plotting the Accuracy and Loss Curves

**Atwanzire Timothy Ian && KATUSHABE MOREEN**

```
In [28]: history
```

```
Out[28]: <keras.callbacks.History at 0x7f2c84420e50>
```

```
In [31]: history.params
```

```
Out[31]: {'verbose': 1, 'epochs': 30, 'steps': 24}
```

```
In [32]: history.history.keys()
```

```
Out[32]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

***loss, accuracy, val loss etc are a python list containing values of loss, accuracy etc at the end of each epoch***

```
In [33]: type(history.history['loss'])
```

```
Out[33]: list
```

```
In [34]: len(history.history['loss'])
```

```
Out[34]: 30
```

```
In [35]: history.history['loss'][:5] # show loss for first 5 epochs
```


```
Out[35]: [3.3086416721343994,  
          2.7993218898773193,  
          2.350374460220337,  
          2.1901614665985107,  
          1.969476342201233]
```

```
In [36]: acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']
```

```
loss = history.history['loss']  
val_loss = history.history['val_loss']
```

```
In [37]: plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(range(EPOCHS), acc, label='Training Accuracy')
plt.plot(range(EPOCHS), val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(range(EPOCHS), loss, label='Training Loss')
plt.plot(range(EPOCHS), val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



Run prediction on a sample image


```
In [38]: import numpy as np
for images_batch, labels_batch in test_ds.take(1):

    first_image = images_batch[0].numpy().astype('uint8')
    first_label = labels_batch[0].numpy()

    print("first image to predict")
    plt.imshow(first_image)
    print("actual label:", class_names[first_label])

    batch_prediction = model.predict(images_batch)
    print("predicted label:", class_names[np.argmax(batch_prediction[0])])
```

```
first image to predict
actual label: blueberries
1/1 [=====] - 2s 2s/step
predicted label: blueberries
```



**Write a function for inference**

```
In [39]: def predict(model, img):  
    img_array = tf.keras.preprocessing.image.img_to_array(images[i].num  
    img_array = tf.expand_dims(img_array, 0)  
  
    predictions = model.predict(img_array)  
  
    predicted_class = class_names[np.argmax(predictions[0])]  
    confidence = round(100 * (np.max(predictions[0])), 2)  
    return predicted_class, confidence
```

**Now run inference on few sample images**



```
In [40]: plt.figure(figsize=(15, 15))
for images, labels in test_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))

        predicted_class, confidence = predict(model, images[i].numpy())
        actual_class = class_names[labels[i]]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class},\n Confidence: {confidence}")

        plt.axis("off")
```

```
1/1 [=====] - 0s 189ms/step
1/1 [=====] - 0s 68ms/step
1/1 [=====] - 0s 83ms/step
1/1 [=====] - 0s 92ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 71ms/step
1/1 [=====] - 0s 75ms/step
1/1 [=====] - 0s 67ms/step
1/1 [=====] - 0s 64ms/step
```

Actual: pineapples,  
Predicted: kiwifruit.  
Confidence: 867.63%



Actual: guava,  
Predicted: grapefruits.  
Confidence: 884.68%



Actual: tomatoes,  
Predicted: apples.  
Confidence: 860.58%



Actual: grapes,  
Predicted: grapes.  
Confidence: 366.15%



Actual: plums,  
Predicted: strawberries.  
Confidence: 1398.9%



Actual: apples,  
Predicted: apples.  
Confidence: 1037.1%



Actual: watermelons,  
Predicted: watermelons.  
Confidence: 978.38%



Actual: blueberries,  
Predicted: blueberries.  
Confidence: 1932.29%



Actual: lemons,  
Predicted: lemons.  
Confidence: 1970.78%





```
In [41]: import os
model_version=max([int(i) for i in os.listdir("./models") + [0]])+1
model.save(f"./models/{model_version}")
```

WARNING:absl:Found untraced functions such as \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_jit\_compiled\_convolution\_op, \_update\_step\_xla while saving (showing 4 of 4). These functions will not be directly callable after loading.

INFO:tensorflow:Assets written to: ./models/4/assets

INFO:tensorflow:Assets written to: ./models/4/assets

```
In [ ]:
```